


In [3]:

```
import numpy as np
```

## CIFAR-10 classifier with CNN


**airplane**




**automobile**



**bird**



**cat**




**deer**




**dog**




**frog**



**horse**



**ship**



**truck**



## Task 1: Forward pass and Backward pass for Convolution

Implement a naive conv\_forward function and apply the vertical and horizontal filter on example images

Implement a conv\_forward using im2col and also apply the filter on example images

Implement a conv\_backward using im2col and do some numerical gradient test

In [1]:

```
from sklearn.datasets import load_sample_image
```

In [4]:

```
china = load_sample_image('china.jpg')[150:220, 130:250] / 255
flower = load_sample_image("flower.jpg")[150:220, 130:250] / 255
images = np.array([china, flower])
batch_size, height, width, channels = images.shape
```

In [5]:


```
import sklearn
import matplotlib.pyplot as plt
```

In [6]:

```
# Create 2 filters
filters = np.zeros(shape=(7, 7, channels, 2), dtype=np.float32)
filters[:, 3, :, 0] = 1 # vertical line
filters[3, :, :, 1] = 1 # horizontal line
```


In [7]:

```
plt.subplot(1, 2, 1)
plt.imshow(filters[:, :, 0, 0])
plt.subplot(1, 2, 2)
plt.imshow(filters[:, :, 0, 1])
plt.show()
```



In [8]:

```
plt.figure(figsize=(15, 15))
plt.subplot(1, 2, 1)
plt.imshow(china)
plt.subplot(1, 2, 2)
plt.imshow(flower)
plt.tight_layout()
plt.show()
```



In [31]:

```
#!pip install torch torchvision -i https://pypi.douban.com/simple
```

In [9]:

```
# if you don't have torch and torchvision  
# uncomment and run the above cell  
import torch  
import torch.nn.functional as F
```

In [12]:

```
filters = np.transpose(filters, (3, 2, 0, 1))  
# We will want our image to be NCHW format for pytorch to process  
images = np.transpose(images, (0, 3, 1, 2))
```

In [13]:

```
# Example on convolution using pytorch  
images_tensor = torch.Tensor(images)  
filters_tensor = torch.Tensor(filters)  
outputs = F.conv2d(images_tensor, filters_tensor, )
```

In [14]:



```
china_convolved = outputs[0].numpy()  
flower_convolved = outputs[1].numpy()
```

In [15]:

```
china_convolved = np.transpose(china_convolved, (1, 2, 0))  
flower_convolved = np.transpose(flower_convolved, (1, 2, 0))
```

In [16]:

```
plt.figure(figsize=(20, 20))  
plt.subplot(1, 2, 1)  
plt.imshow(china_convolved[:, :, 0], cmap='gray')  
plt.subplot(1, 2, 2)  
plt.imshow(china_convolved[:, :, 1], cmap='gray')  
plt.show()
```



In [ ]:

```
def conv_forward_naive(x, w, b=None, pad=0, strides=1):
    """
    A naive implementation of the forward pass for a convolutional layer.
    The input consists of N data points, each with C channels, height H and width W. We convolve each input with F different filters, where each filter spans all C channels and has height HH and width WW.
    Input:
    - x: Input data of shape (N, C, H, W)
    - w: Filter weights of shape (F, C, HH, WW)
    - b: Biases, of shape (F,)
    - stride: The number of pixels between adjacent receptive fields in the horizontal and vertical directions.
    - pad: The number of pixels that will be used to zero-pad the input.
    Returns a tuple of:
    - out: Output data, of shape (N, F, H', W') where H' and W' are given by
         $H' = 1 + (H + 2 * pad - HH) / stride$ 
         $W' = 1 + (W + 2 * pad - WW) / stride$ 
    - cache: (x, w, b, pad, strides)
    """
    out = None
    N, C, H, W = x.shape
    F, C, HH, WW = w.shape
    S = strides
    P = pad

    if b is None:
        b = np.zeros((F))

    """
    TODO: write your code here
    """

    cache = (x, w, b, pad, strides)
    return out, cache
```

In [35]:

```
"""
    Here apply you implmentation of convolution on the sample images
    and show the result
"""
```

Out[35]:

```
'\n    Here apply you implmentation of convolution on the sample images\n    and s\nhow the result\n'
```

## im2col



it turns out that we can use a clever algorithm called im2col to convert a image to a matrix and the convolution (which we naively implement using for-loops) will be a single matrix multiplication

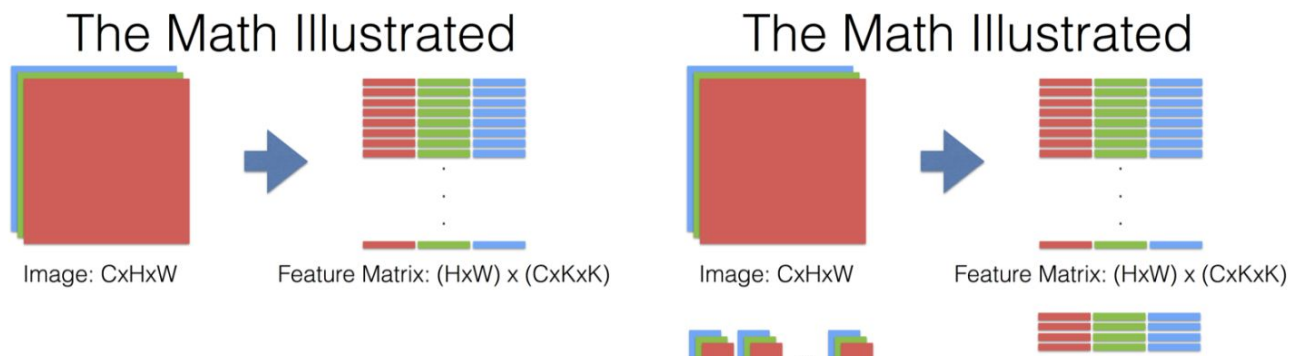
check this [url \(http://cs231n.github.io/convolutional-networks/#conv\)](http://cs231n.github.io/convolutional-networks/#conv) and Yangqing's memo [Conv in Caffe \(https://github.com/Yangqing/caffe/wiki/Convolution-in-Caffe:-a-memo\)](https://github.com/Yangqing/caffe/wiki/Convolution-in-Caffe:-a-memo).

# Optimizing Conv in Caffe

- Solution: Converting convolution to GEMM  
This is done in two (lazy) steps:  
(1) do an im2col to convert the image to a matrix  
(2) call GEMM to do the actual computation
- Pro: simple logic, easy implementation
- Con: LOTS of memory overhead  
(and it does not utilize the flops/param advantage very well)

\* More rants at Yangqing's memo: [Convolution in Caffe](https://github.com/Yangqing/caffe/wiki/Convolution-in-Caffe:-a-memo)





In [17]:

```
def im2col(image, ksize, stride):
    """
    TODO: Implement a im2col to NCHW format images
    """
    pass
```

In [18]:

```
def conv_forward_im2col(x, w, b=None, pad=0, strides=1):
    """
    TODO: Implement conv_forward using im2col
    """
    pass
```

In [36]:

```
"""
Here apply you implmentation of convolution on the sample images
and show the result
"""
```

Out[36]:

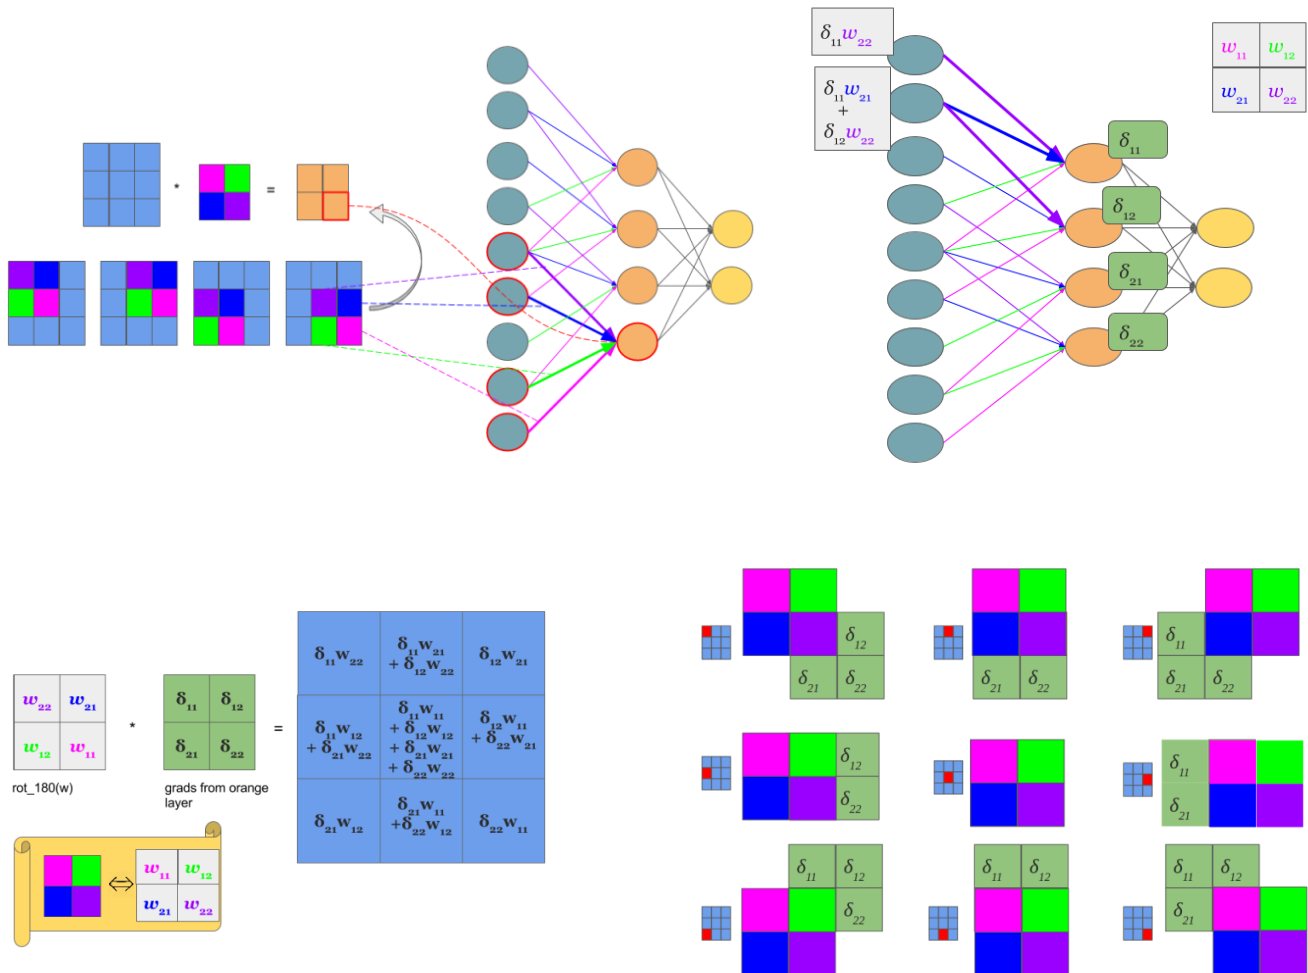
```
'\n Here apply you implmentation of convolution on the sample images\n and s
how the result\n'
```

In [19]:

```
def conv_backward_im2col(dout, cache):
    """
    TODO
    """
    pass
```

We can actually see that the gradient through the convolution is the rotated kernel convolved over the error from later layer

(if you prefer solid math over this visual proof, check [this](https://grzegorzwardys.wordpress.com/2016/04/22/8/) (<https://grzegorzwardys.wordpress.com/2016/04/22/8/>) [this](https://www.jefkine.com/general/2016/09/05/backpropagation-in-convolutional-neural-networks/) (<https://www.jefkine.com/general/2016/09/05/backpropagation-in-convolutional-neural-networks/>) and this lecture notes ([http://courses.cs.tau.ac.il/Caffe\\_workshop/Bootcamp/pdf\\_lectures/Lecture%203%20CNN%20-%20backpropagation.pdf](http://courses.cs.tau.ac.il/Caffe_workshop/Bootcamp/pdf_lectures/Lecture%203%20CNN%20-%20backpropagation.pdf)))



In [20]:

```
def eval_numerical_gradient_array(f, x, df, h=1e-5):
    """
    Evaluate a numeric gradient for a function that accepts a numpy
    array and returns a numpy array.
    """
    grad = np.zeros_like(x)
    it = np.nditer(x, flags=['multi_index'], op_flags=['readwrite'])
    while not it.finished:
        ix = it.multi_index

        oldval = x[ix]
        x[ix] = oldval + h
        pos = f(x).copy()
        x[ix] = oldval - h
        neg = f(x).copy()
        x[ix] = oldval

        grad[ix] = np.sum((pos - neg) * df) / (2 * h)
        it.iternext()
    return grad
```

You relative error from the below cell output should be around 1e-9

In [ ]:

```
x = np.random.randn(4, 3, 5, 5)
w = np.random.randn(2, 3, 3, 3)
b = np.random.randn(2,)
dout = np.random.randn(4, 2, 5, 5)
stride, pad = 1, 1

dx_num = eval_numerical_gradient_array(lambda x: conv_forward_im2col(x, w, b, pad, stride)[0],
x, dout)
dw_num = eval_numerical_gradient_array(lambda w: conv_forward_im2col(x, w, b, pad, stride)[0],
w, dout)
db_num = eval_numerical_gradient_array(lambda b: conv_forward_im2col(x, w, b, pad, stride)[0],
b, dout)

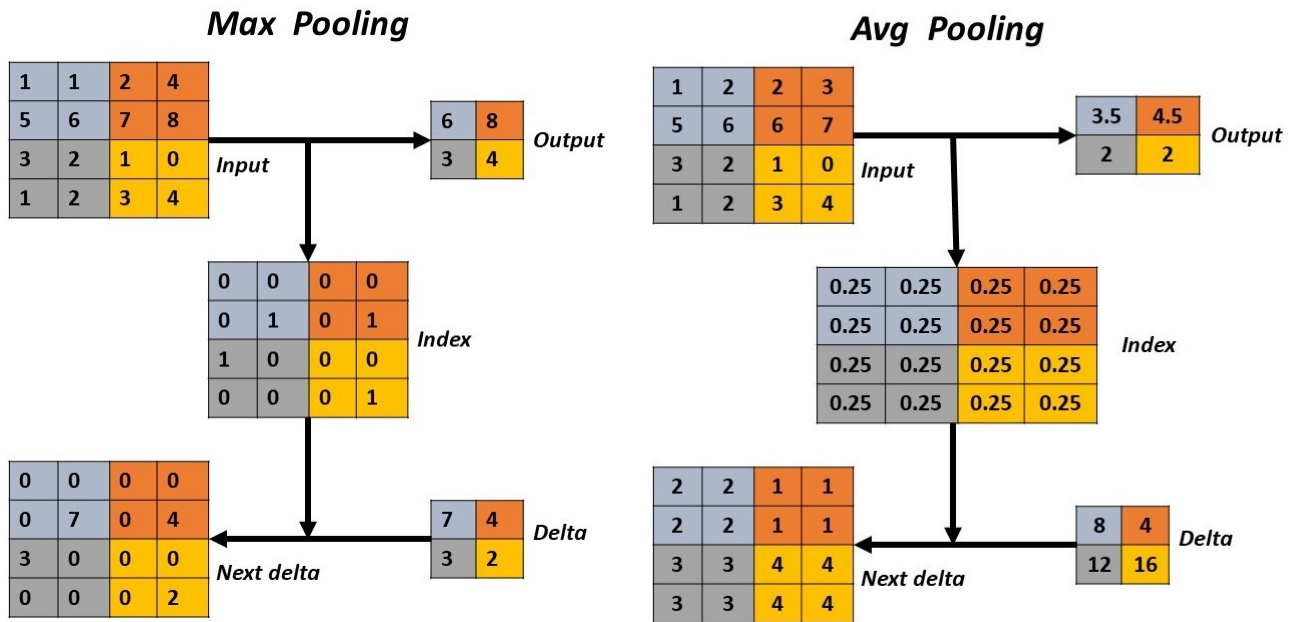
out, cache = conv_forward_im2col(x, w, b, pad, stride)
dx, dw, db = conv_backward_im2col(dout, cache)

def rel_error(x, y):
    """ returns relative error """
    return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))

print("dx error: ", rel_error(dx, dx_num))
print("dw error: ", rel_error(dw, dw_num))
print("db error: ", rel_error(db, db_num))
```



## Task 2: Forward pass and backward pass of other layers



you can choose either pool method and implement

In [22]:

```
def pool_forward(x, height, width, ksize, method):
    """
    TODO
    """
    pass

def pool_backward(x, height, width, ksize, method):
    """
    TODO
    """
    pass
```

In [23]:

```
def fc_forward(x, w, b):
    """
    TODO
    """
    pass

def fc_backward(dout, cache):
    """
    TODO
    """
    pass
```

In [24]:

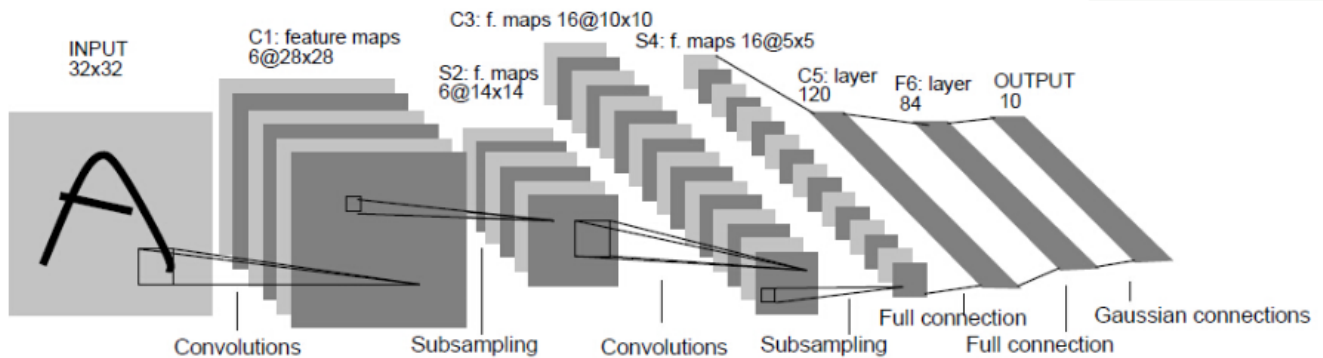
```
def relu_forward(x):  
    """  
    TODO  
    """  
    pass  
def relu_backward(dout, cache):  
    """  
    TODO  
    """  
    pass
```

In [26]:

```
def softmax_loss(x, y):  
    """  
    Computes the loss and gradient for softmax classification.  
  
    Inputs:  
    - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class  
      for the ith input.  
    - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and  
      0 <= y[i] < C  
  
    Returns a tuple of:  
    - loss: Scalar giving the loss  
    - probs: The predicted probability  
    - dx: Gradient of the loss with respect to x  
    """  
    """  
    TODO  
    """  
    pass
```

## Task3: Combine all this together and train a CNN on CIFAR-10!

You can use an arch similar to LeNet



1. overfit a small data(achieve 100% accuracy on a small data)
2. train on the whole training data, validation on validation data, finally test on test data

In [30]:

```
#!/pip install torch torchvision -i https://pypi.douban.com/simple
```

In [29]:

```
import torchvision

# I'm using torchvision to load the CIFAR10 dataset
# if you don't have torch and torchvision
# uncomment and run the above cell
def load_cifar(path='./data'):
    cifar_train = torchvision.datasets.cifar.CIFAR10(path, download=True)
    cifar_test = torchvision.datasets.cifar.CIFAR10(path, download=True, train=False)

    cifar_train_img_list, cifar_train_label_list = [], []
    for train_example in cifar_train:
        img = np.array(train_example[0])[np.newaxis, :]
        label = train_example[1]
        cifar_train_img_list.append(img)
        cifar_train_label_list.append(label)
    cifar_train_img = np.concatenate(cifar_train_img_list, axis=0).astype(np.float) / 255
    cifar_train_label = np.array(cifar_train_label_list).astype(np.int)

    cifar_test_img_list, cifar_test_label_list = [], []
    for test_example in cifar_test:
        img = np.array(test_example[0])[np.newaxis, :]
        label = test_example[1]
        cifar_test_img_list.append(img)
        cifar_test_label_list.append(label)
    cifar_test_img = np.concatenate(cifar_test_img_list, axis=0).astype(np.float) / 255
    cifar_test_label = np.array(cifar_test_label_list).astype(np.int)

    return cifar_train_img[:49000], cifar_train_label[:49000], \
           cifar_train_img[49000:], cifar_train_label[49000:], \
           cifar_test_img, cifar_test_label
```

In [32]:

```
X, y, X_val, y_val, X_test, y_test = load_cifar()
X = np.transpose(X, (0, 3, 1, 2))
X_val = np.transpose(X_val, (0, 3, 1, 2))
X_test = np.transpose(X_test, (0, 3, 1, 2))
batch_size = 64
X_mini_batches = [X[k:k+batch_size] for k in range(0, len(X), batch_size)]
y_mini_batches = [y[k:k+batch_size] for k in range(0, len(y), batch_size)]
X_dummy = X[:100]
y_dummy = y[:100]
```

Files already downloaded and verified  
Files already downloaded and verified

In [33]:

```
X.shape, y.shape, X_val.shape, y_val.shape, X_test.shape, y_test.shape
```

Out[33]:

```
((49000, 3, 32, 32),
 (49000,),
 (1000, 3, 32, 32),
 (1000,),
 (10000, 3, 32, 32),
 (10000,))
```

## Check if you have any bug by overfitting a small dat

In [34]:

```
X_dummy.shape, y_dummy.shape
```

Out[34]:

```
((100, 3, 32, 32), (100,))
```

In [ ]:

```
for epoch in range(200):
    alpha = 0.1
    reg = 0.01
    # forward pass through network using X_dummy
    """
        TODO
    """

    #from ipdb import set_trace; set_trace()
    # Calculate the loss
    """
        TODO
    """

    acc = np.mean(np.argmax(pred, axis=1) == y_dummy)
    if epoch % 20 == 0:
        #print('dloss: ', dloss)
        print('epoch:', epoch, 'loss:', loss)
        print('epoch:', epoch, 'reg_loss:', reg_loss)
        print('epoch:', epoch, 'acc:', acc)

    # backward pass through network
    """
        TODO
    """

    # update parameter
    """
        TODO
    """
```

## Train on the whole dataset

```
for epoch in range(num_epoch):
    for X_batch, y_batch in zip(X_batches, y_batches):
        forward(X_batch)
        loss(pred, y_batch)
        backward()
        update()
    if epoch % interval == 0:
        validate(X_val, y_val)

test(X_test, y_test)
```

In [ ]: